

2. Deployment Strategy (Production-Ready)

2.1 Containerization

- **Dockerization of Microservices:**
 - Each microservice has its own Dockerfile, ensuring separation of concerns.
 - Containers encapsulate all dependencies and runtime configurations.
 - **Multi-Stage Builds:**
 - Use **multi-stage Docker builds** to keep images small and secure.
 - Final image built on a **distroless base image**, minimizing attack surface and removing unnecessary tools (e.g., bash, curl).
 - **Best Practices:**
 - Use .dockerignore to prevent leaking sensitive files.
 - Regularly scan images using tools like **Trivy** or **Grype**.
-

2.2 Orchestration

- **Development Environment:**
 - Use **Docker Compose** to spin up multiple services locally.
 - Supports volume mounting, service dependencies, and .env injection.
 - **Production Environment:**
 - Use **Kubernetes (K8s)** for orchestration.
 - Deployment manifests located under deploy/k8s/
 - Supports horizontal scaling, self-healing, service discovery, and resource limits.
 - **Compatible Platforms:**
 - Can be deployed on **Amazon EKS, Google GKE, Azure AKS**, or self-hosted K8s clusters.
 - **Namespace Management:**
 - Use namespaces to isolate environments (e.g., dev, staging, prod).
-

2.3 CI/CD

- **CI/CD Tools:**
 - Use **GitHub Actions**, **GitLab CI**, or **CircleCI** for automation.
 - **Pipeline Stages:**
 1. **Linting:** Enforce code style using golangci-lint
 2. **Testing:** Run unit and integration tests
 3. **Build:** Compile and build Docker images
 4. **Push:** Push versioned images to Docker Hub / ECR / GCR
 5. **Deploy:** Roll out changes to Kubernetes using Helm or Kustomize
 - **Deployment Strategies:**
 - Support for **Blue/Green** or **Canary Deployments** using **ArgoCD** or **Flux**.
 - Enable automatic rollback on health check failure.
-

2.4 Secrets Management

- **Storage Options:**
 - **Kubernetes Secrets:** For environment-specific secrets mounted at runtime.
 - **AWS Secrets Manager** or **HashiCorp Vault:** For dynamic and secure secret injection.
 - **Security Measures:**
 - Never commit `.env` or `.secret` files to version control.
 - Rotate secrets periodically and enforce least privilege access.
-

3. Non-Functional Requirements (NFRs)

3.1 Logging

- **Structured Logging:**
 - Use **uber-go/zap** or **rs/zerolog** for structured logs.
 - Logs are in **JSON format** for compatibility with aggregators.
- **Log Aggregation:**

- Logs are collected using **FluentBit** and forwarded to:
 - **ElasticSearch** (via EFK stack: Elasticsearch, FluentBit, Kibana)
 - **Grafana Loki** for a lightweight, cost-effective solution
- **Example Log Entry:**

```
{  
  "level": "info",  
  "service": "auth",  
  "msg": "user login successful",  
  "userID": "123",  
  "timestamp": "2025-08-03T12:00:00Z"  
}
```

3.2 Authentication & Authorization

- **Authentication:**
 - Use **JWTs** for access and refresh tokens
 - Middleware extracts and validates token, injecting claims into request context
 - **Authorization:**
 - Use **Role-Based Access Control (RBAC)**:
 - Roles: Admin, User, Vendor
 - Configurable via database or environment variables
 - **Future Integration Options:**
 - Support for **OAuth2 / OpenID Connect** via **Google, Auth0, Keycloak**
-

3.3 Monitoring & Observability

- **Metrics Collection:**
 - Use **Prometheus** to scrape service metrics
 - Dashboards built with **Grafana**

- Key Metrics:
 - HTTP request count & duration
 - DB query performance
 - Memory/CPU usage per container
 - **Distributed Tracing:**
 - Use **OpenTelemetry (OTEL)** SDKs for tracing
 - Backend options: **Jaeger, Tempo**
 - **Health Probes:**
 - Implement /healthz endpoints for:
 - **Readiness probe:** Determines if app is ready to serve
 - **Liveness probe:** Determines if app should be restarted
-

3.4 Rate Limiting & Throttling

- **Enforcement Points:**
 - API Gateway (e.g., **Kong, Nginx Ingress**) applies global rate limits
 - Per-user or per-IP rate limiting via middleware
 - **Implementation:**
 - Use **Redis-backed** token bucket using libs like **ulule/limiter**
 - **Benefits:**
 - Prevents DDoS and abuse
 - Ensures fair usage across users
-

3.5 Caching

- **Cache Layer:**
 - Use **Redis** for:
 - Session and token storage
 - Frequently queried data (e.g., restaurant list, menu items)
- **Caching Strategy:**

- Set **TTL (Time-to-Live)** on keys
 - Manual invalidation on updates (e.g., restaurant updated)
 - **Optional:** In-memory LRU caching for performance-sensitive reads
-

3.6 Secure Communication

- **Ingress Security:**
 - TLS termination at **Ingress Controller** (e.g., Nginx, Traefik)
 - Use **Let's Encrypt** or **cert-manager** for auto-renewed certs
 - **Internal Security:**
 - gRPC/HTTP traffic restricted via **Kubernetes Network Policies**
 - Enable **mTLS** for service-to-service communication if required
-

3.7 Database & Storage

- **Transactional Database:**
 - **MySQL** used for persistent relational data
 - Ensure regular **automated backups** using MySQL dump, Percona XtraBackup, or managed snapshots (e.g., AWS RDS)
 - Tune queries and indexes for high throughput under load
- **Migrations:**
 - Use **Golang-migrate** or **goose** for version-controlled schema changes
 - Run migrations as part of CI/CD pipeline before deployment
- **Object Storage:**
 - Use **AWS S3, MinIO**, or compatible service
 - For user uploads (e.g., images, documents)
 - Ensure access via signed URLs and configure proper IAM policies